

Open Elective Course [OE]

Course Code: CSO507

Winter 2023-24

Lecture#

Deep Learning

Unit-1: Machine Learning Basics [Part-III]


Course Instructor:

Dr. Monidipa Das

Assistant Professor

Department of Computer Science and Engineering

Indian Institute of Technology (Indian School of Mines) Dhanbad, Jharkhand 826004, India




Optimization




- **Optimization** is concerned with optimizing an **objective function**
 - finding the value of an argument that minimizes or maximizes the function
 - Most optimization algorithms are formulated in terms of minimizing a function $f(x)$
 - Maximization is accomplished via minimizing the negative of an objective function (e.g., minimize $-f(x)$)
 - In minimization problems, the objective function is often referred to as a **cost function** or **loss function** or **error function**
- **Optimization is very important for machine learning**
 - The performance of optimization algorithms affects the model's training efficiency

Convex and Non-Convex Functions



A concave function:
no line segment joining
two points on the graph
lies above the graph
at any point




A convex function:
no line segment joining
two points on the graph
lies below the graph
at any point





A function that is neither
concave nor convex:
the line segment shown lies
above the graph at some
points and below it at others


Negative of a convex function is
called a **concave** function, which
also has a unique optima (maxima)



Convex functions are bowl-shaped. They have a unique
optima (minima)




Non-convex functions have multiple minima. Usually harder
to optimize as compared to convex functions




Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Convex Functions



- In mathematical terms, the function f is a **convex function** if for all points x_1, x_2 and for all $\lambda \in [0,1]$:

$$\lambda f(x_1) + (1 - \lambda)f(x_2) \geq f(\lambda x_1 + (1 - \lambda)x_2)$$





Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Functions and their optima




- Many ML problems require us to optimize a function f of some variable(s) x
- For simplicity, assume f is a scalar-valued function of a scalar x ($f: \mathbb{R} \rightarrow \mathbb{R}$)





- Any function has one/more optima (maxima, minima), and maybe saddle points
- Finding the optima or saddle points requires derivatives/gradients of the function

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Saddle Points




- Points where derivative is zero but are neither minima nor maxima




- Saddle points are very common for loss functions of deep learning models
 - The optimization algorithms may stall at saddle points, without reaching a minima
 - Need to be handled carefully during optimization
- Second or higher derivative may help identify if a stationary point is a saddle

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Derivatives




- How the derivative itself changes tells us about the function's optima




- The second derivative $f''(x)$ can provide this information

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Stationary/Critical Points



- Are the points where the derivative of the function (say $f(x)$) is zero, i.e., $f'(x) = 0$
- The stationary points can be:
 - Minimum**, a point where the derivative changes from negative to positive
 - Maximum**, a point where the derivative changes from positive to negative
 - Saddle point**, derivative is either positive or negative on both sides of the point
- The minimum and maximum points are collectively known as **extremum points**
- The nature of stationary points can be determined based on the second derivative of $f(x)$ at the point
 - If $f''(x) > 0$, the point is a minimum
 - If $f''(x) < 0$, the point is a maximum
 - If $f''(x) = 0$, inconclusive, the point can be a saddle point, but it may not



P, Q, R are called Stationary Points

The same concept also applies to gradients of multivariate functions


Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad



Methods for Solving Optimization Problems


Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Method 1: Non-iterative & Using First-Order Optimality



- First order optimality: The gradient \mathbf{g} must be equal to zero at the optima

$$\mathbf{g} = \nabla_{\mathbf{w}}[L(\mathbf{w})] = \mathbf{0}$$




Called “first order” since only gradient is used and gradient provides the first order info about the function being optimized

The approach works only for very simple problems where the objective is convex and there are no constraints on the values \mathbf{w} can take

- Sometimes, setting $\mathbf{g} = \mathbf{0}$ and solving for \mathbf{w} gives a closed form solution
- If closed form solution is not available, the gradient vector \mathbf{g} can still be used in iterative optimization algos, like [gradient descent](#)

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad


Method 2: Iterative Optimization via Gradient Descent



- Iterative since it requires several steps/iterations to find the optimal solution
- Suppose $y = f(x)$, x, y real nos.
 - Derivative of function denoted: $f'(x)$ or as $\frac{dy}{dx}$
 - Derivative $f'(x)$ gives the slope of $f(x)$ at point x
 - It specifies how to scale a small change in input to obtain a corresponding change in the output: $f(x + \epsilon) \approx f(x) + \epsilon f'(x)$
 - It tells how you make a small change in input to make a small improvement in y
 - We know that $f(x - \epsilon \text{sign}(f'(x)))$ is less than $f(x)$ for small ϵ .
 - Thus we can reduce $f(x)$ by moving x in small steps with opposite sign of derivative
 - This technique is called **gradient descent (Cauchy 1847)**


Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

How Gradient Descent uses derivatives




- Criterion $f(x)$ minimized by moving from current solution in direction of the negative of gradient

- For $x > 0$, $f(x)$ increases with x and $f'(x) > 0$
- For $x < 0$, $f(x)$ decreases with x and $f'(x) < 0$
- Use $f'(x)$ to follow function downhill
- Reduce $f(x)$ by going in direction opposite sign of derivative $f'(x)$



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad



Method of Gradient Descent




- The gradient points directly uphill, and the negative gradient points directly downhill
- Thus we can decrease f by moving in the direction of the negative gradient
 - This is known as the method of **steepest descent** or **gradient descent**
- Steepest descent proposes a new point $\mathbf{x}' = \mathbf{x} - \alpha \nabla_{\mathbf{x}} f(\mathbf{x})$
 - where α is the **learning rate**, a positive scalar. Set to a small constant.
- Steepest descent converges when every element of the gradient is zero
 - In practice, very close to zero
- We may be able to avoid iterative algorithm and jump to the critical point by solving the equation $\nabla_{\mathbf{x}} f(\mathbf{x}) = 0$ for \mathbf{x}

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Learning rate is very important




Too large, the next point will perpetually bounce haphazardly across the bottom of the well



Too small learning rate will take too long

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient descent algorithm



Gradient Descent: Algorithmic representation

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ (simultaneously update $j = 0$ and $j = 1$)
 }
 Parameters Cost function, which is to be minimized

Correct updating of parameters


temp0 := $\theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$
 temp1 := $\theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$
 $\theta_0 :=$ temp0
 $\theta_1 :=$ temp1

Incorrect updating of parameters

temp0 := $\theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$
 $\theta_0 :=$ temp0
 temp1 := $\theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$
 $\theta_1 :=$ temp1

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent for Linear Regression




$$h_{\theta}(x) = f_{\theta}(x) = \theta_0 + \theta_1 x \quad J(\theta_0, \theta_1) = \frac{1}{2N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

m is the number of samples/ data points

repeat until convergence {
 $\theta_0 := \theta_0 - \alpha \frac{1}{N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})$
 $\theta_1 := \theta_1 - \alpha \frac{1}{N} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$
 }
 Update θ_0 and θ_1 simultaneously

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution




Example Regression Problem:

Prediction of house price in a city based on living area.

Living area (feet ²)	Price (1000\$)
2104	400
1600	330
2400	369
1416	232
3000	540
\vdots	\vdots


Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

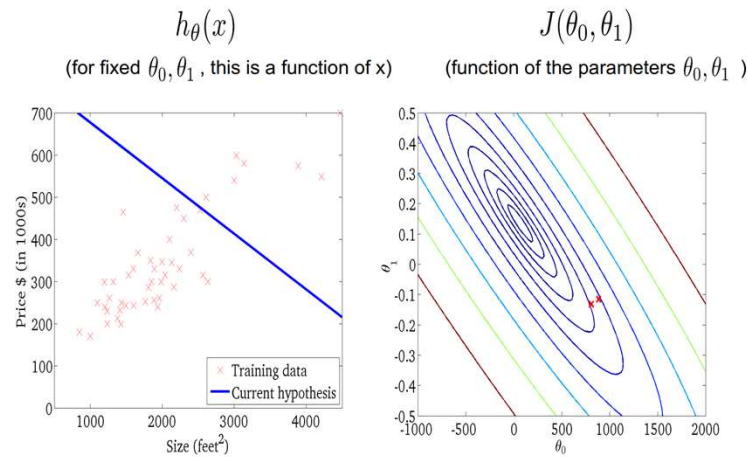

(function of the parameters θ_0, θ_1)



Courtesy: Andrew Ng

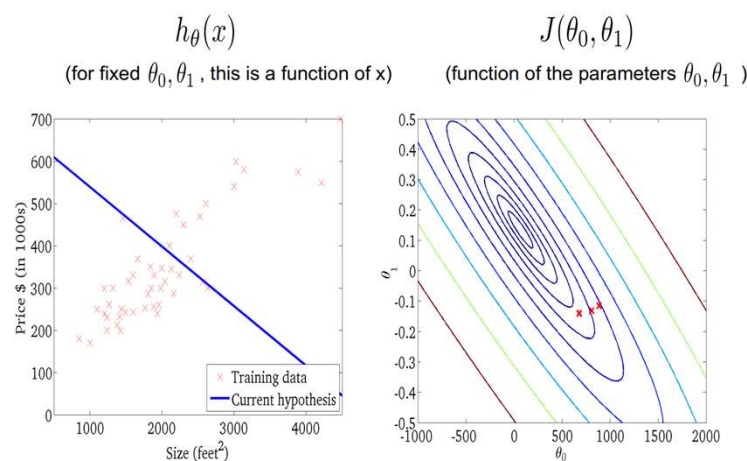

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution



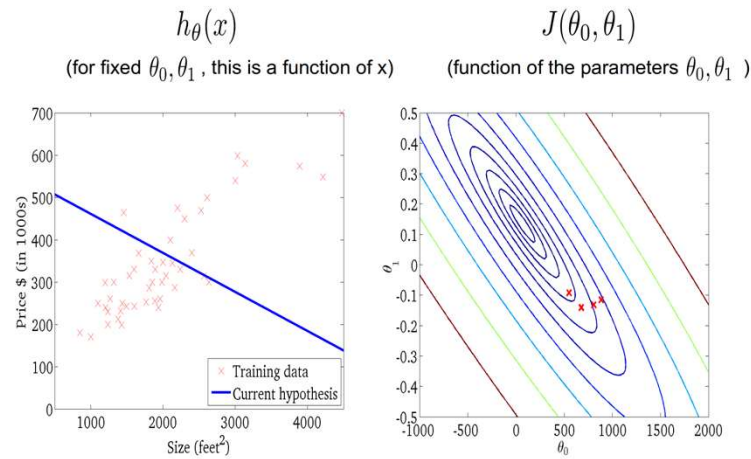

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution



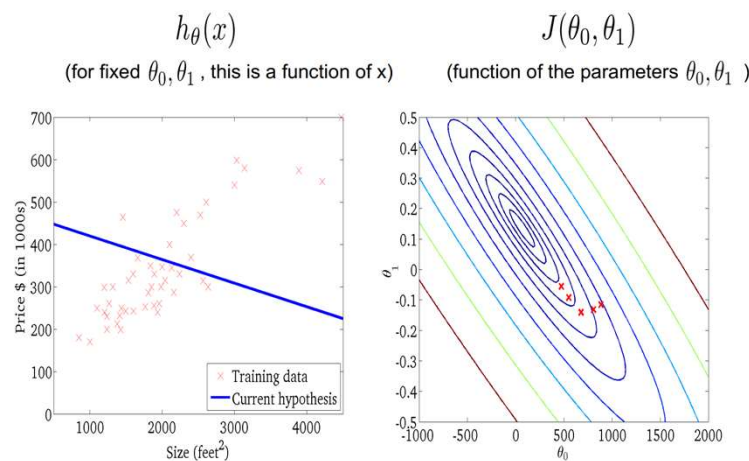

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution



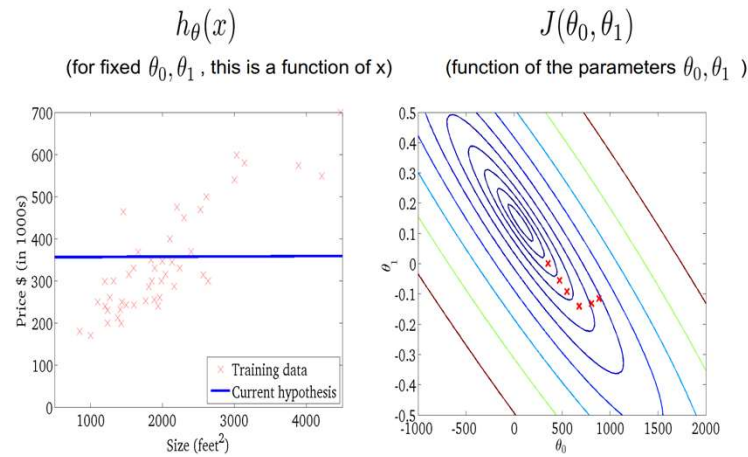

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution



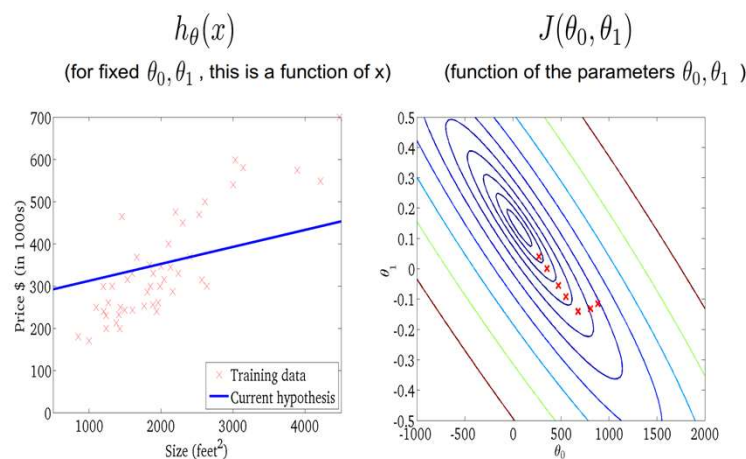

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution



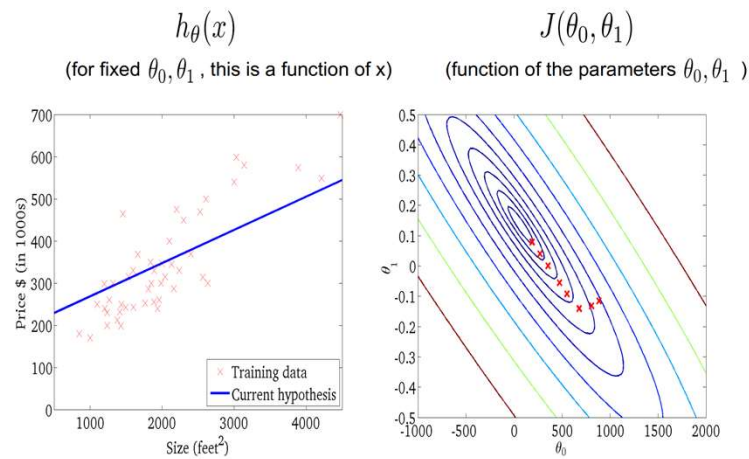

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution



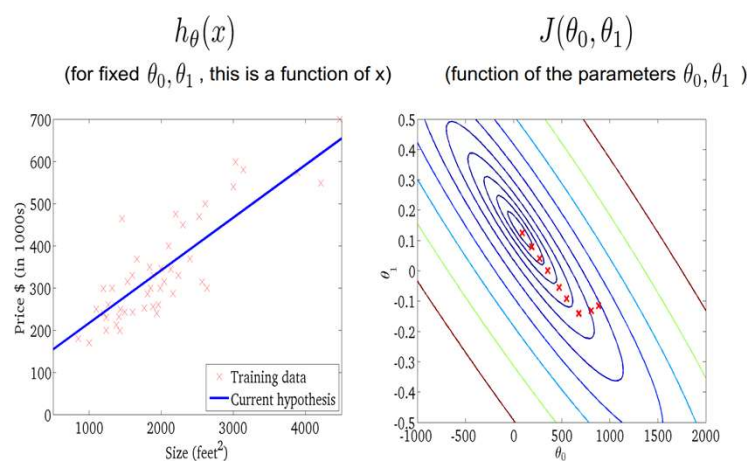

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution




Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent in Execution




Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

GD: Advantages and Disadvantages




- Advantages
 - Ease of implementation
 - Convergence guarantee
 - Scalability
 - Versatility
- Disadvantages
 - Sensitivity to learning rate
 - Sensitivity to initialization
 - Convergence speed
 - Correctness




Courtesy: Andrew Ng

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent: Facts (Summary)



- Gradient gives the direction of **steepest change** in function's value
- Iterative since it requires several steps/iterations to find the optimal solution
- For convex functions, GD will converge to the global minima
- The gradient descent algorithm is not written for all types of functions.
- The function has to satisfy two conditions for Gradient Descent to be applicable on it:
 - differentiable
 - convex function




- Good initialization needed for non-convex functions
- The learning rate is very imp. Should be set carefully (fixed or chosen adaptively).
- For max. problems we can use gradient **ascent**. Will move in the direction of the gradient

$$w^{(t+1)} = w^{(t)} + \alpha_t g^{(t)} \leftarrow \text{Gradient}$$

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Types of Gradient Descent



- There are **three major types** of Gradient Descent Algorithms:
 - Batch Gradient Descent (BGD)
 - Stochastic Gradient Descent (SGD)
 - Mini-Batch Gradient Descent (MBGD)

Gradient Descent (GD): Pseudo-Code


Input: parameters (θ), gradient of the loss function with respect to the parameters ($d\theta$), learning rate (α)

Update parameters: $\theta = \theta - \alpha \times d\theta$

Output: updated parameters (θ)

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Pseudo-Code: Batch Gradient Descent (BGD)




Step 1: Select a learning rate α

Step 2: Select initial parameter values θ as the starting point

Step 3: Update all parameters from the gradient of the training data set, i.e. compute


$$\theta = \theta - \alpha \times \nabla_{\theta} J(\theta)$$

Step 4: Repeat Step 3 until a local minima is reached




Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Pseudo-Code: Stochastic Gradient Descent (SGD)



- Step 1: Randomly shuffle the data set of size N
- Step 2: Select a learning rate α
- Step 3: Select initial parameter values θ as the starting point
- Step 4: Update all parameters from the gradient of a single training example x^j, y^j , i.e. compute
- $$\theta = \theta - \alpha \times \nabla_{\theta} J(\theta; x^j; y^j)$$
- Step 5: Repeat Step 4 until a local minimum is reached


Stochastic Gradient Descent



Common to decay learning rate linearly until iteration τ : $\alpha_i = (1-\epsilon)\alpha_0 + \epsilon\alpha_{\tau}$ with $\epsilon = i/\tau$
 After iteration τ , it is common to leave α constant. Often a small positive value in the range 0.0 to 1.0

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Pseudo-Code: Mini-Batch Gradient Descent (MBGD)



- Estimate gradient on a batch of b samples ($b < N$), N is the no. of training samples

Step 1: Randomly shuffle the data set of size N


Step 2: Select a learning rate α

Step 3: Select initial parameter values θ as the starting point

Step 4: Update all parameters from the gradient calculated against a batch size of b training examples, i.e. compute


$$\theta = \theta - \alpha \times \nabla_{\theta} J(\theta; x^{j:j+b}; y^{j:j+b})$$

Step 5: Repeat Step 4 until a local minimum is reached



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad


Momentum Method




- SGD is a popular optimization strategy but it can be slow
- Momentum method accelerates learning, when:
 - Facing high curvature
 - Small but consistent gradients
 - Noisy gradients
- It works by accumulating the moving average of past gradients and moves in that direction while exponentially decaying

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Gradient Descent with Momentum




- Gradient descent with momentum converges faster than standard gradient descent
- Taking large steps in w_2 direction and small steps in w_1 direction slows down algorithm



- Momentum reduces oscillation in w_2 direction
- Now can set a higher learning rate

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Momentum Definition



- Introduce velocity variable v
- This is the direction and speed at which parameters move through parameter space
- Name momentum comes from physics and is mass times velocity
 - The momentum algorithm assumes unit mass
- A hyperparameter $\delta \in [0,1)$ determines exponential decay of v

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Momentum Update Rule



- The update rule is given by

$$\begin{aligned} v &\leftarrow \delta v - \alpha \nabla_{\theta} \left(\frac{1}{m} \sum_{i=1}^m L(f(x^{(i)}; \theta), y^{(i)}) \right) \\ \theta &\leftarrow \theta + v \end{aligned}$$


- The velocity v accumulates the gradient elements

$$\nabla_{\theta} \left(\frac{1}{m} \sum_{i=1}^m L(f(x^{(i)}; \theta), y^{(i)}) \right)$$

- The larger δ is relative to α , the more previous gradients affect the current direction

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

SGD Algorithm with Momentum



```


INPUT: cost function  $J(\theta)$ , learning rate  $\alpha$ , number of iterations:  $T$ ,
      batch size  $B$ , initial velocity:  $v$ 
INITIALIZE: random  $\theta$ 
FOR  $i = 1$  to  $T$  DO
  Split the training examples into  $B$  mini-batches of size  $b$ :
  FOR  $j = 1$  to number of mini-batches  $B$  DO
    Compute the gradient of  $J$  with respect to  $\theta$  for a mini-batch
    of training examples:
     $\text{gradient} = 1/b \times \nabla_{\theta} \sum (J(\theta, x^j, y^j))$ 
    Compute velocity update:  $v = \delta v - \alpha \times \text{gradient}$ 

    Update the parameters  $\theta$ :
     $\theta = \theta + v$ 
  END FOR
END FOR
OUTPUT:  $\theta$ 

```

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Nesterov Momentum



- A variant to accelerate gradient, with update


$$v \leftarrow \delta v - \alpha \nabla_{\theta} \left[\frac{1}{m} \sum_{i=1}^m L(f(x^{(i)}; \theta + \delta v), y^{(i)}) \right],$$

$$\theta \leftarrow \theta + v,$$

- where parameters δ and α play a similar role as in the standard momentum method
- Difference between Nesterov and standard momentum is where gradient is evaluated.
 - Nesterov gradient is evaluated after the current velocity is applied.
 - One can interpret Nesterov as attempting to add a correction factor to the standard method of momentum

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

SGD with Nesterov Momentum




- A variant of the momentum algorithm
 - Nesterov's accelerated gradient method
- Applies a correction factor to standard method

```

INPUT: cost function  $J(\theta)$ , learning rate  $\alpha$ , number of iterations:  $T$ ,
      batch size  $B$ , initial velocity:  $v$ 
INITIALIZE: random  $\theta$ 
FOR  $i = 1$  to  $T$  DO
  Split the training examples into  $B$  mini-batches of size  $b$ :
  FOR  $j = 1$  to number of mini-batches  $B$  DO
    Apply interim update:  $\tilde{\theta} = \theta + \delta v$ 
    Compute the gradient of  $J$  with respect to  $\theta$  for a mini-batch
    of training examples:
    gradient =  $1/b \times \nabla_{\tilde{\theta}} \Sigma(J(\tilde{\theta}, x^j, y^j))$ 
    Compute velocity update:  $v = \delta v - \alpha \times \text{gradient}$ 
    Update the parameters  $\theta$ :
     $\theta = \theta + v$ 
  END FOR
END FOR
OUTPUT:  $\theta$ 

```

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad



Questions?

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad