

Open Elective Course [OE]

Course Code: CSO507

Winter 2023-24

Lecture#

Deep Learning

Unit-8: Generative Models (Part-III)


Course Instructor:

Dr. Monidipa Das

Assistant Professor

Department of Computer Science and Engineering

Indian Institute of Technology (Indian School of Mines) Dhanbad, Jharkhand 826004, India



Variational Autoencoders



- ① We modelled the autoencoding process using a probabilistic (Bayesian) framework, with an observed and a hidden variable
- ② We wanted to calculate the posterior distribution $p_{\theta}(z|x)$, but this is complicated
- ③ We used an approximation $q_{\phi}(z|x)$ to $p_{\theta}(z|x)$
- ④ We used the ELBO as a loss function to minimise $KL(q_{\phi}(z|x)||p_{\theta}(z))$
 - Ensures a good reconstruction
 - Encourages the latent space to follow our chosen distribution (the prior $p_{\theta}(z)$)


Variational Autoencoders



- Probabilistic spin to traditional autoencoders => allows generating data
- Defines an intractable density => derive and optimize a (variational) lower bound
- **Advantages**
 - Theoretically-motivated, loss function meaningful
 - Learn to and from mapping (encoder and decoder)
- **Drawbacks**
 - In practice, samples blurrier and lower quality
 - Have to re-write loss function for each different model, not always easy

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Variational Autoencoders



VAEs define intractable density function with latent variables z (that we need to marginalize):

$$p_{\theta}(x) = \int p_{\theta}(z)p_{\theta}(x|z)dz$$

cannot optimize directly, derive and optimize lower bound of likelihood instead

What if we give up on explicitly modeling density, and just want to sample?

GANs: don't work with any explicit density function

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Generative Adversarial Networks (GANs)




Sample from a simple distributions, e.g., random noise. Learn transformation to the training distribution

Question: What can we use to represent complex transformation function?


Input: Random noise

Output: Sample from training distribution



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Generative Adversarial Networks (GANs)



- **Generative**
 - Learn a generative model
- **Adversarial**
 - Trained in an adversarial setting
- **Networks**
 - Use Deep Neural Networks

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

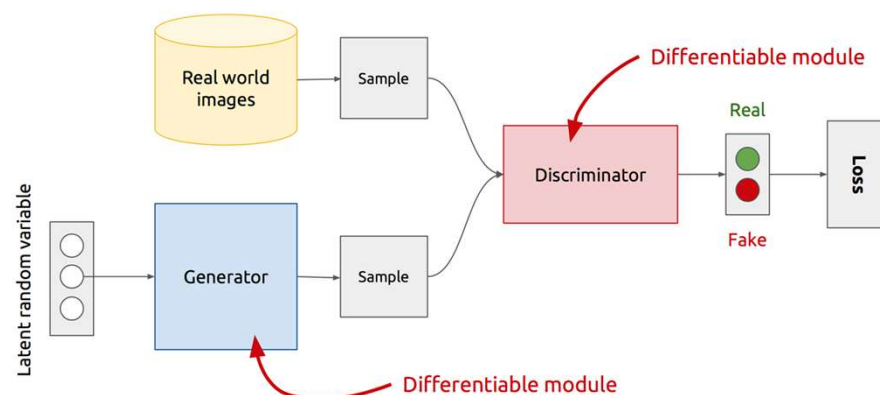

Generative Adversarial Networks (GANs)



- Adversarial Training
 - Can generate adversarial samples to fool a discriminative model
 - Can use those adversarial samples to make models robust
 - Require more effort to generate adversarial samples
 - Repeat this and we get better discriminative model

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

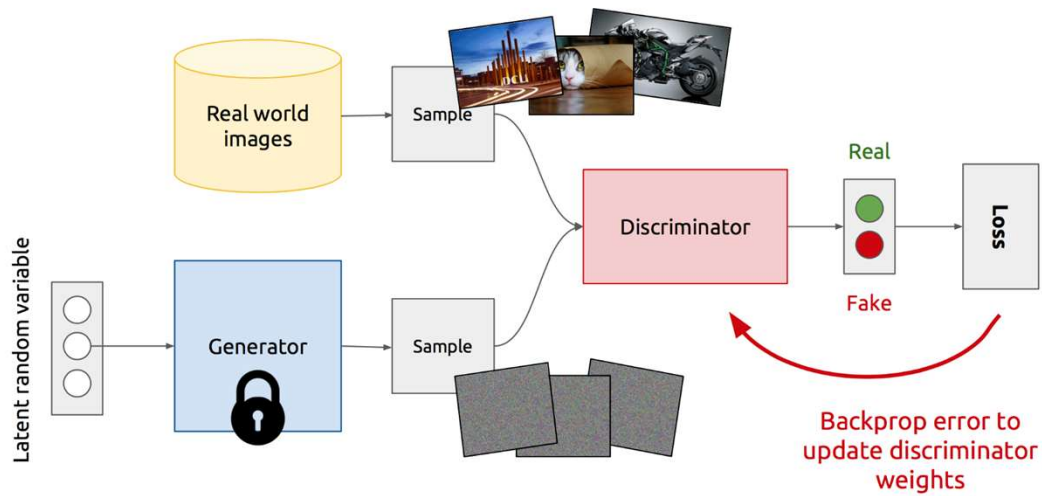
GAN's Architecture



- Z is some random noise (Gaussian/Uniform).
- Z can be thought as the latent representation of the image.

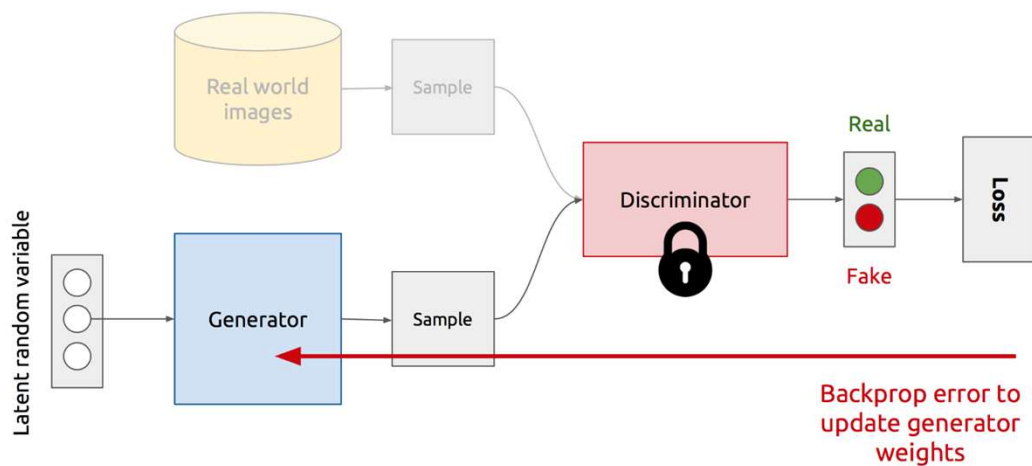

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Training Discriminator




Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Training Generator



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Training GAN: Formulation



- Training GNN: Two-player Game

$$\min_G \max_D V(D, G)$$

- Formulated as a **minimax game**, where:


- The Discriminator is trying to maximize its reward $V(D, G)$
- The Generator is trying to minimize Discriminator's reward (or maximize its loss)

$$V(D, G) = \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Discriminator output
for real data x
Discriminator output for
generated fake data G(z)

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Training GAN



Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. Gradient **ascent** on discriminator


$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. Gradient **descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Training GAN



Discriminator
updates

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log(1 - D(G(z^{(i)}))) \right].$$

end for

Generator
updates

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:


$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(1 - D(G(z^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Training GAN



Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. Gradient **ascent** on discriminator


$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{\text{data}}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. Gradient **descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$


Gradient signal
dominated by region
where sample is
already good

When sample is likely
fake, want to learn from
it to improve generator.



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Training GAN



Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$


Alternate between:

1. Gradient **ascent** on discriminator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$


2. Instead, gradient **ascent** on generator, different objective

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

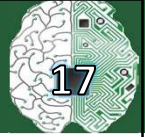
Advantages of GAN



- Plenty of existing work on Deep Generative Models but GANs
 - Don't take explicit density function
 - Game theoretic Approach
 - Better Sample generation
- Why GANs?
 - Sampling (or generation) is straightforward
 - Training doesn't involve Maximum Likelihood estimation
 - Robust to Overfitting

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad


Disadvantages of GAN



- **Probability Distribution is Implicit**
 - Not straightforward to compute $P(X)$.
 - Thus Vanilla GANs are only good for Sampling/Generation
- **Training is Hard**
 - **Non-Convergence**
 - SGD is not designed to find equilibrium
 - **Mode-Collapse**
 - Can focus on a few realistic images from the training dataset

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Non-Convergence



- **Deep Learning models (in general) involve a single player**
 - The player tries to maximize its reward (minimize its loss).
 - Use SGD (with Backpropagation) to find the optimal parameters.
 - SGD has convergence guarantees (under certain conditions).
 - **Problem:** With non-convexity, we might converge to local optima.

$$\min_G L(G)$$


- **GANs instead involve two (or more) players**
 - Discriminator is trying to maximize its reward.
 - Generator is trying to minimize Discriminator's reward.

$$\min_G \max_D V(D, G)$$


- SGD was not designed to find the Nash equilibrium of a game.
- **Problem:** We might not converge to the Nash equilibrium at all.

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Mode-Collapse




- Generator fails to output diverse samples



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Generative Adversarial Network



Optimal GAN discriminator D^*

For a fixed G , the optimal D is $D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$


$$\begin{aligned}
 \mathcal{L}(G, D) &= \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \\
 &= \int_{\mathcal{X}} p_{data}(x) \log(D(x)) dx + \int_{\mathcal{Z}} p_z(z) \log(1 - D(G(z))) dz \\
 &= \int_{\mathcal{X}} p_{data}(x) \log(D(x)) + p_G(x) \log(1 - D(x)) dx.
 \end{aligned}$$

- For every x , the maximum of the previous equation w.r.t $D(x)$ is

$$D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$$

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Generative Adversarial Network



Global optimum of the GAN loss function

The global optimum of the GAN loss function is achieved if and only if $p_G = p_{data}$. At this point $\mathcal{L}(G, D) = -\log 4$.

- First, our previous lemma allows us to rewrite the loss function


$$\max_D \mathcal{L}(G, D) = \mathbb{E}_{x \sim p_{data}} [\log D^*(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D^*(G(z)))] \quad (3)$$

$$\begin{aligned} &= \mathbb{E}_{x \sim p_{data}} [\log D^*(x)] + \mathbb{E}_{x \sim p_G} [\log(1 - D^*(x))] \\ &= \mathbb{E}_{x \sim p_{data}} \left[\log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + \mathbb{E}_{x \sim p_G} \left[\log \left(\frac{p_G(x)}{p_{data}(x) + p_G(x)} \right) \right]. \end{aligned} \quad (4)$$

- Therefore, if $p_G = p_{data}$, then $\mathcal{L}(G, D) = \log \frac{1}{2} + \log \frac{1}{2} = -\log(4)$

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Generative Adversarial Network



- Now, we are going to show that $-\log(4)$ is the optimal value of the loss function
- First, we remark that :


$$-\log(4) = \mathbb{E}_{x \sim p_{data}} [-\log(2)] + \mathbb{E}_{x \sim p_G} [-\log(2)]. \quad (5)$$

- Therefore, by subtracting Equation 5 from Equation 4, we have

$$\begin{aligned} \mathcal{L}(G, D^*) &= -\log(4) + \int p_{data}(x) \log \frac{p_{data}(x)}{\frac{1}{2}(p_{data}(x) + p_G(x))} dx + \\ &\quad \int p_G(x) \log \frac{p_G(x)}{\frac{1}{2}(p_{data}(x) + p_G(x))} dx \\ &= -\log(4) + KL \left(p_{data} \parallel \frac{p_{data} + p_G}{2} \right) + KL \left(p_G \parallel \frac{p_{data} + p_G}{2} \right). \end{aligned}$$

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Generative Adversarial Network



- This can also be rewritten as

$$\mathcal{L}(G, D^*) = -\log(4) + 2 \text{JSD}(\mathbf{p}_{\text{data}} \parallel \mathbf{p}_G).$$


- The **JSD** is the **Jensen-Shannon divergence**
- This is another distance between distributions. For p and q , we have :

$$\text{JSD}(p, q) = \frac{1}{2}KL\left(p \parallel \frac{1}{2}(p + q)\right) + \frac{1}{2}KL\left(q \parallel \frac{1}{2}(p + q)\right)$$

- The JSD is **non-negative and equal to zero if and only if**
 $p_{\text{data}} = p_G$
- Therefore $-\log(4)$ is the optimal value, and only reached when
 $p_{\text{data}} = p_G$

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

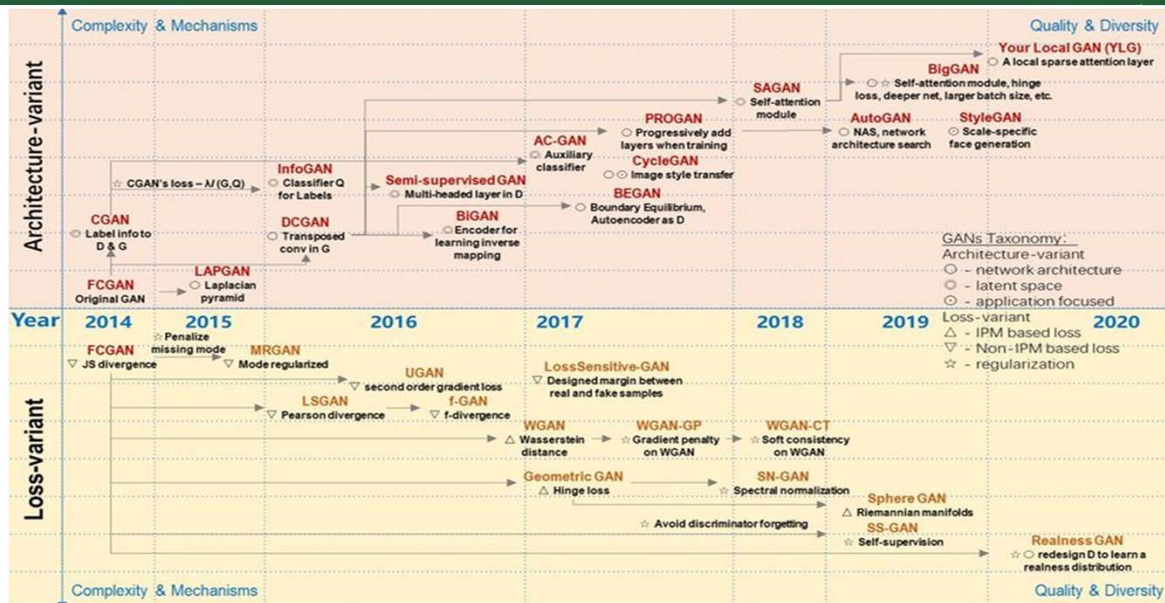

GAN Variants



- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BIGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Evolution of Generative Model



Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Laplacian Pyramid of Adversarial Networks



- Based on the Laplacian Pyramid representation of images. (1983)
- Generate high resolution (dimension) images by using a hierarchical system of GANs
- Iteratively increase image resolution and quality.

Denton, E.L., Chintala, S. and Fergus, R., 2015. "Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks". NIPS (2015)

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Laplacian Pyramid of Adversarial Networks

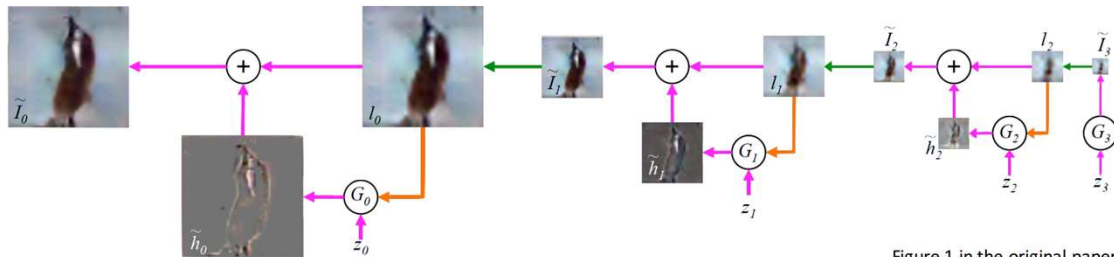


Figure 1 in the original paper.

Image Generation using a LAPGAN

- Generator G_3 generates the base image \tilde{I}_3 from random noise input z_3 .
- Generators (G_2, G_1, G_0) iteratively generate the *difference image* (\tilde{h}) **conditioned on previous small image (I)**.
- This *difference image* is added to an **up-scaled version of previous smaller image**.

Denton, E.L., Chintala, S. and Fergus, R., 2015. "Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks". NIPS (2015)

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Laplacian Pyramid of Adversarial Networks

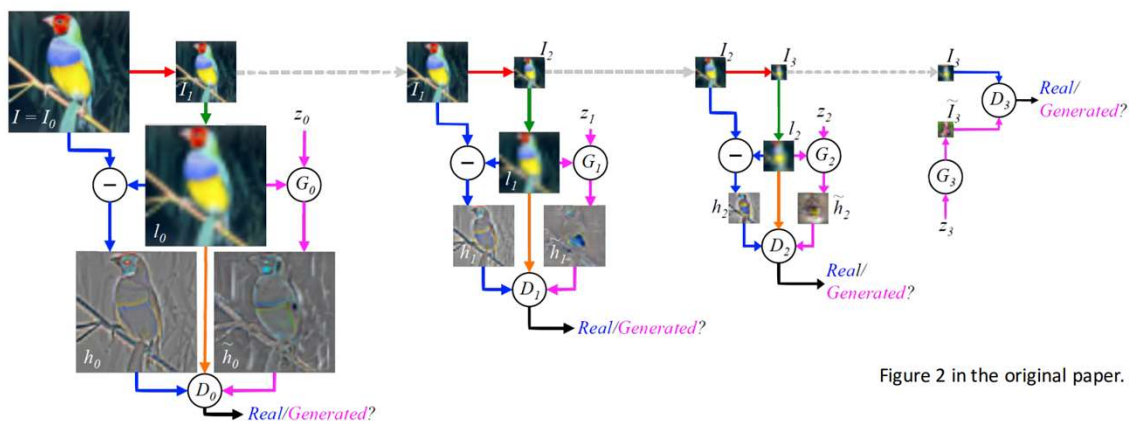


Figure 2 in the original paper.

Training Procedure:

Models at each level are trained independently to learn the required representation.

Denton, E.L., Chintala, S. and Fergus, R., 2015. "Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks". NIPS (2015)

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Coupled GAN






Figure 2 in the original paper.

- Learning a *joint distribution* of *multi-domain* images.
- Using GANs to learn the joint distribution with samples drawn from the marginal distributions.
- Direct applications in domain adaptation and image translation.

Liu, Ming-Yu, and Onel Tuzel. "Coupled generative adversarial networks". NIPS (2016).

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Coupled GANs



- Architecture




Figure 1 of the original paper.

Liu, Ming-Yu, and Onel Tuzel. "Coupled generative adversarial networks". NIPS (2016).

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

Coupled GANs

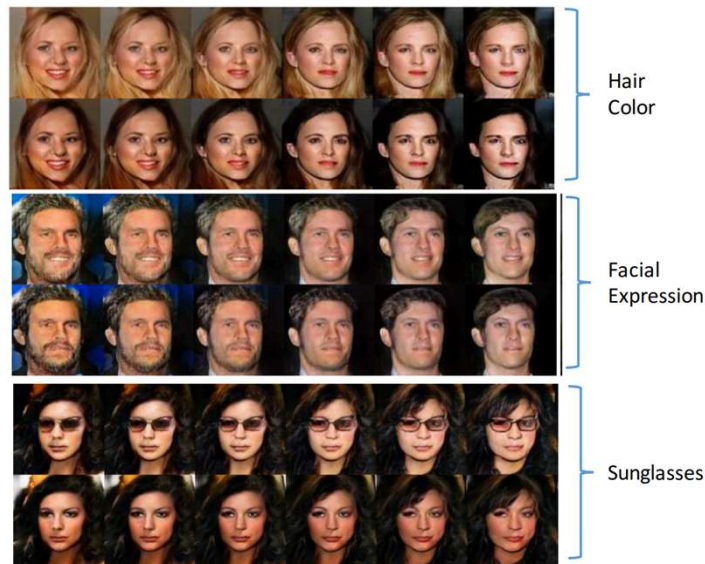




Figure 4 in the original paper.

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

StyleGAN Explained



StyleGAN attempts to tackle the limitation of GAN by adding progressive training to adjust each detail level separately. In doing so, **user can control visual features expressed in each detail level in an isolated manner without affecting other levels.**

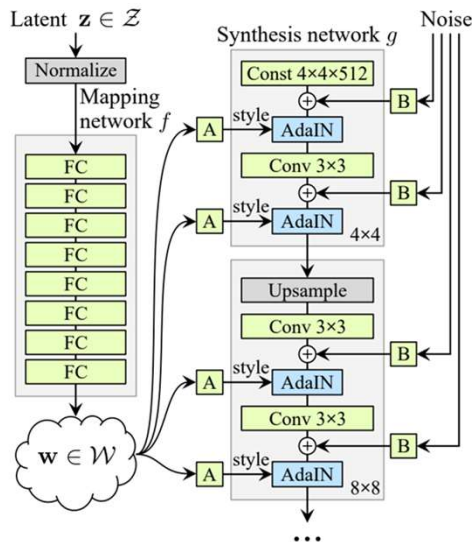

| Detail level | Resolution | What is affected? |
|--------------|--------------|--|
| Coarse | Up to 82 | Pose, general hair style, face shape etc. |
| Middle | 162 to 322 | Finer facial features, hair style, eyes open/closed etc. |
| Fine | 642 to 10242 | Color scheme (eye, hair and skin) and micro features. |

Our generator thinks of an image as a collection of “styles”, where each style controls the effects at a particular scale

- Coarse styles → pose, hair, face shape
- Middle styles → facial features, eyes
- Fine styles → color scheme

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad

StyleGAN Model Architecture



StyleGAN is described as a **progressive growing** GAN architecture with 5 modifications, each of which was added and evaluated incrementally:

Baseline Progressive GAN

Addition of tuning and bilinear up-sampling


Addition of mapping network and AdaIN (styles)

Removal of latent vector input to generator

Addition of noise to each block.

Addition of Mixing regularization

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad



Questions?

Prof. Monidipa Das, Department of CSE, IIT (ISM) Dhanbad